



The University of Texas at Austin  
Oden Institute for Computational  
Engineering and Sciences

April 29, 2022

# Satellite Image Classification

---

CSE 382M – Foundations of Data Science  
Final Project

**Julie Pham**

PhD Student – Aerospace Engineering & Engineering Mechanics

**Sebastian Henao-Garcia**

PhD Student – Computational Science, Engineering, and Mathematics

**Vignesh Sella**

PhD Student – Computational Science, Engineering, and Mathematics

# Outline

---

1. Introduction to the dataset and the problem
2. Dimensionality reduction of the feature space
3. Implementation of classifiers
  1. k-Nearest Neighbors (k-NN)
  2. Support Vector Machine (SVM)
  3. Convolutional Neural Network (CNN)
4. Comparison of the classifiers

# Image Classification: Dataset

---

Remote Sensing Image Scene Classification (NWPU-RESISC45)<sup>[1]</sup> image dataset

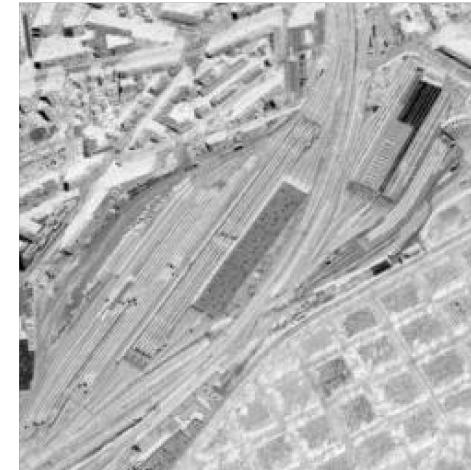
- 256 x 256 pixel resolution, spatial resolution ~30 m to 20 cm
- Subset selection of 4 classes with 700 images each
- 60% training, 15% validation, 25% testing



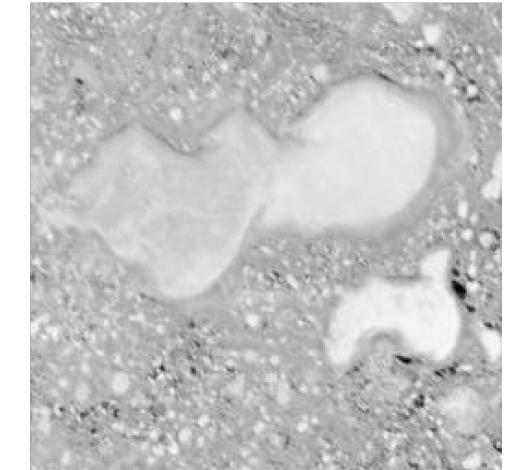
Basketball court



Baseball diamond

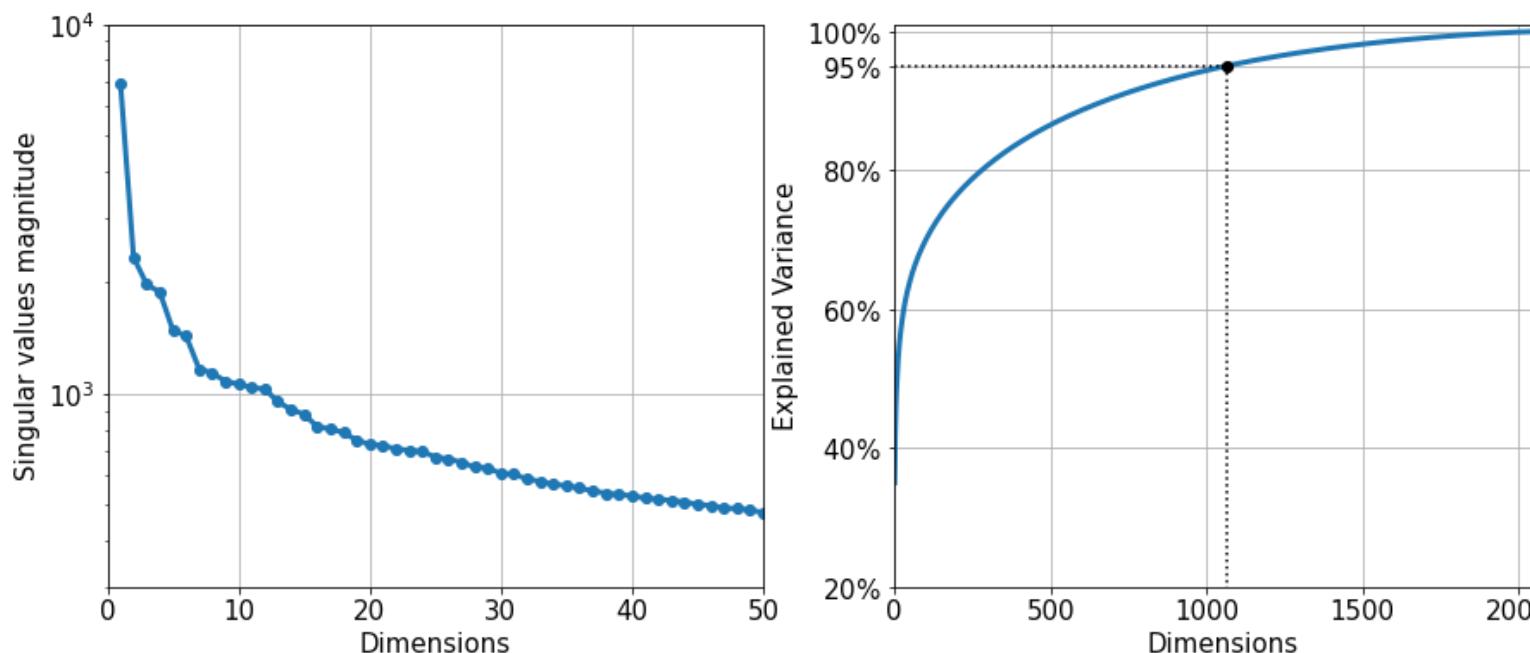
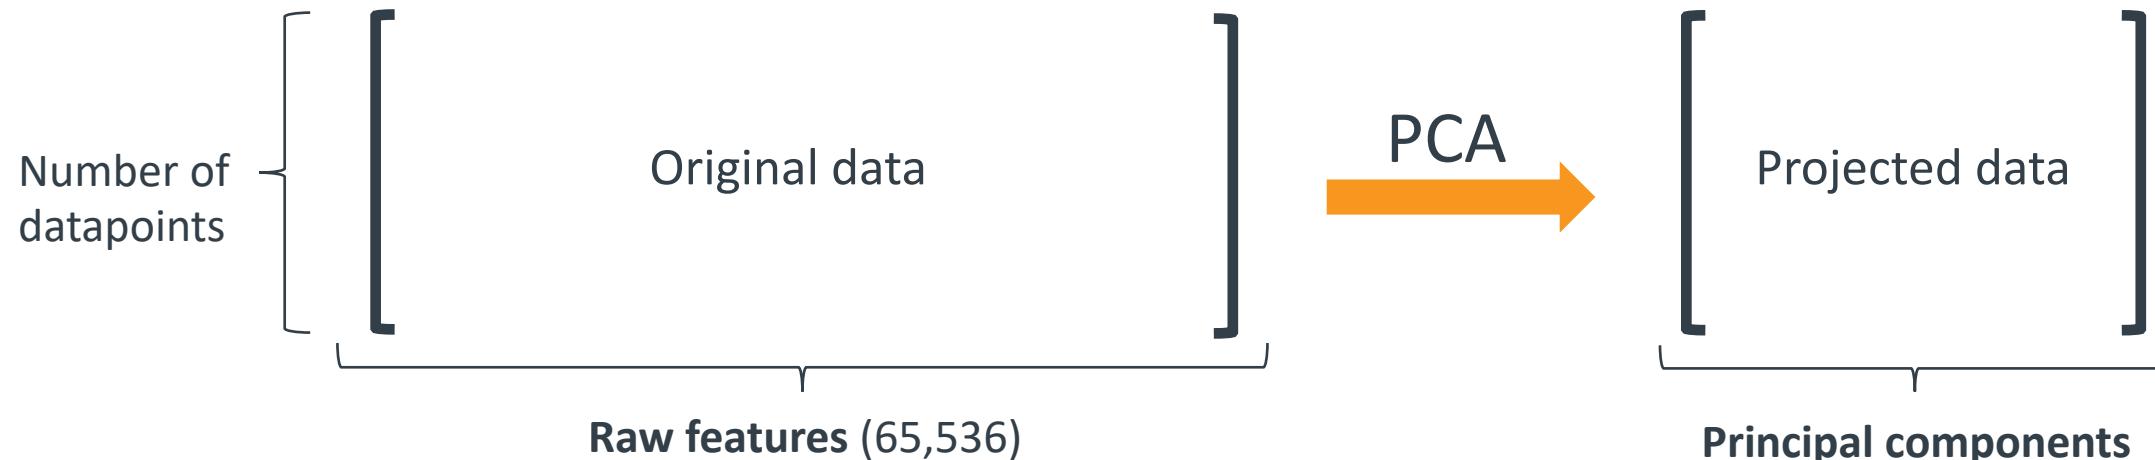


Railway station



Lake

# Dimensionality Reduction via PCA



Principal components:  
 $XV = U\Sigma$

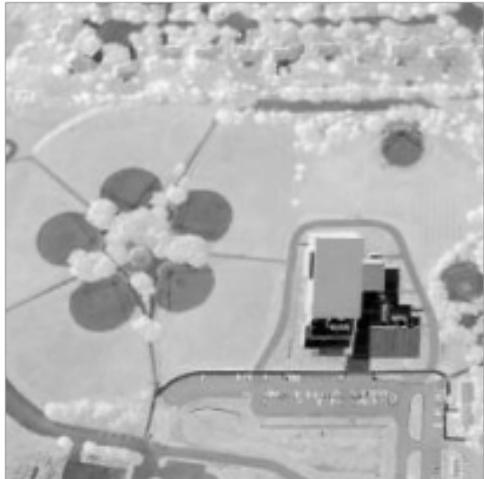
Explained variance  
$$\frac{\sum_{i=1}^k \sigma_i^2}{\sum_{i=1}^N \sigma_i^2}$$

# Dimensionality Reduction via PCA

---

Training data

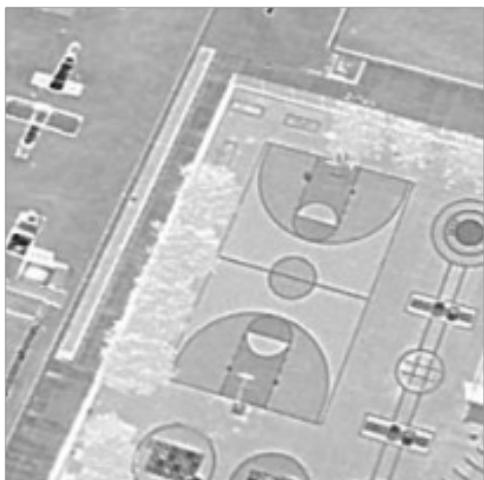
Original



Reduced



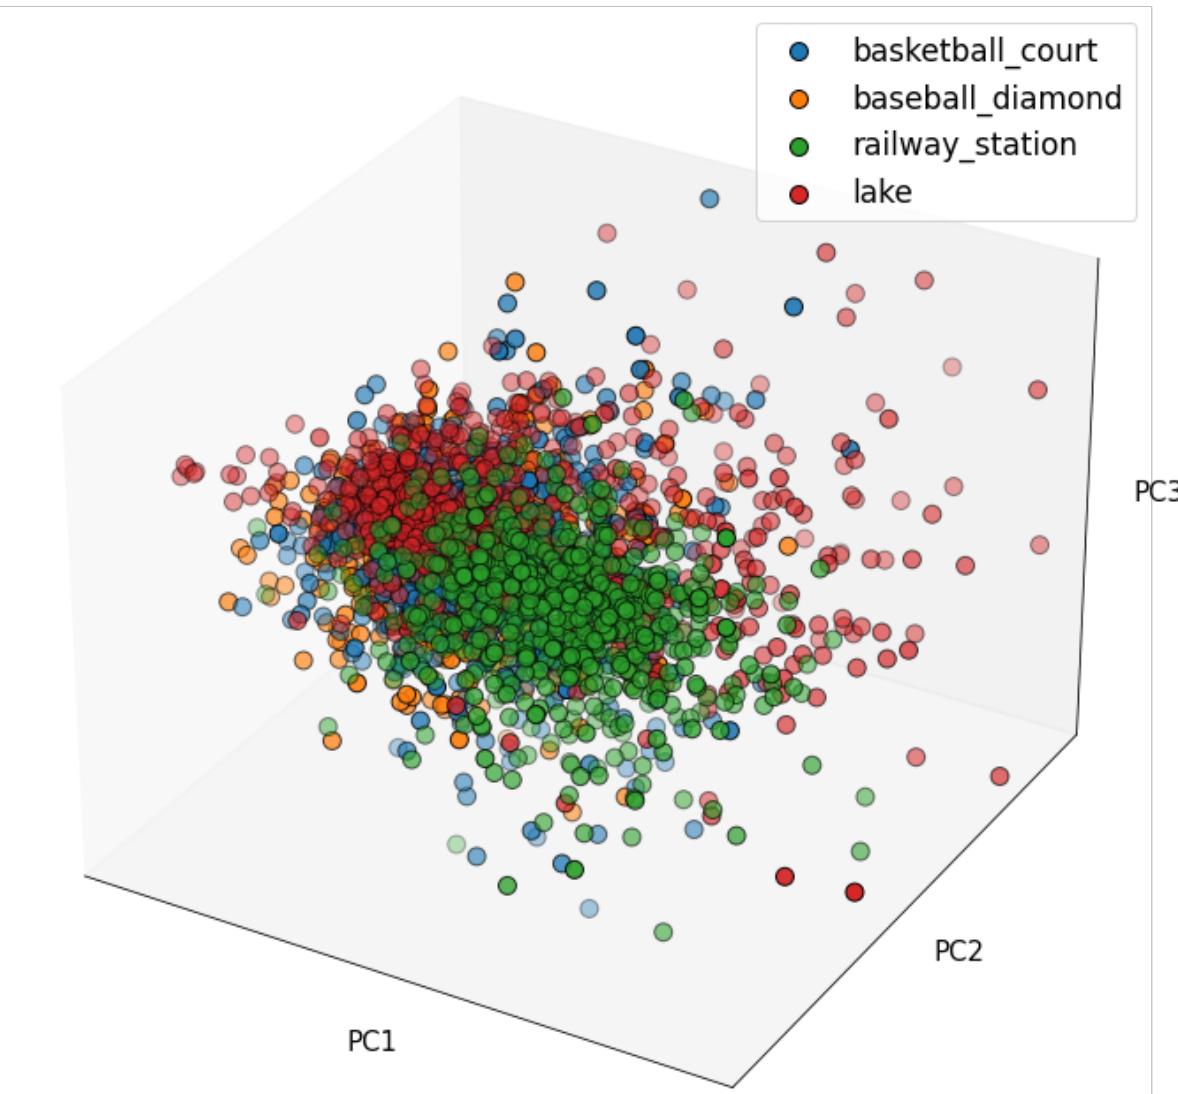
Test data



The PCA modes are computed using the training data set only

The test data set is then projected onto the computed PCA modes

# Dimensionality Reduction via PCA



Using only 3 reduced features, we see some separation in the data, but it is still not strong enough to draw a clear decision boundary.

Then, we need to include more features, but the optimal number depends on the nature of the classification algorithm to be trained.

# Our goal

---

To train three different machine learning models to predict the class a given image belongs to based on:

k Nearest Neighbors  
(KNN)

Support Vector Machine  
(SVM)

Convolutional Neural  
Network  
(CNN)

Comparing them through:

Prediction accuracy

Misclassifications  
among classes

# k-NN algorithm

---

1. Compute (Euclidean) distance between query point and all points in training dataset
2. Select k points with the smallest distances to query point
3. Obtain labels corresponding to the k points
4. Count how many of each label there are
5. Predict the mode

# k-NN: Curse of Dimensionality

---

What happens when your data is high-dimensional?  
...the Euclidean distances look equally dissimilar!

Dimension reduction is necessary in this case,  
but centering/scaling does not make a big difference.

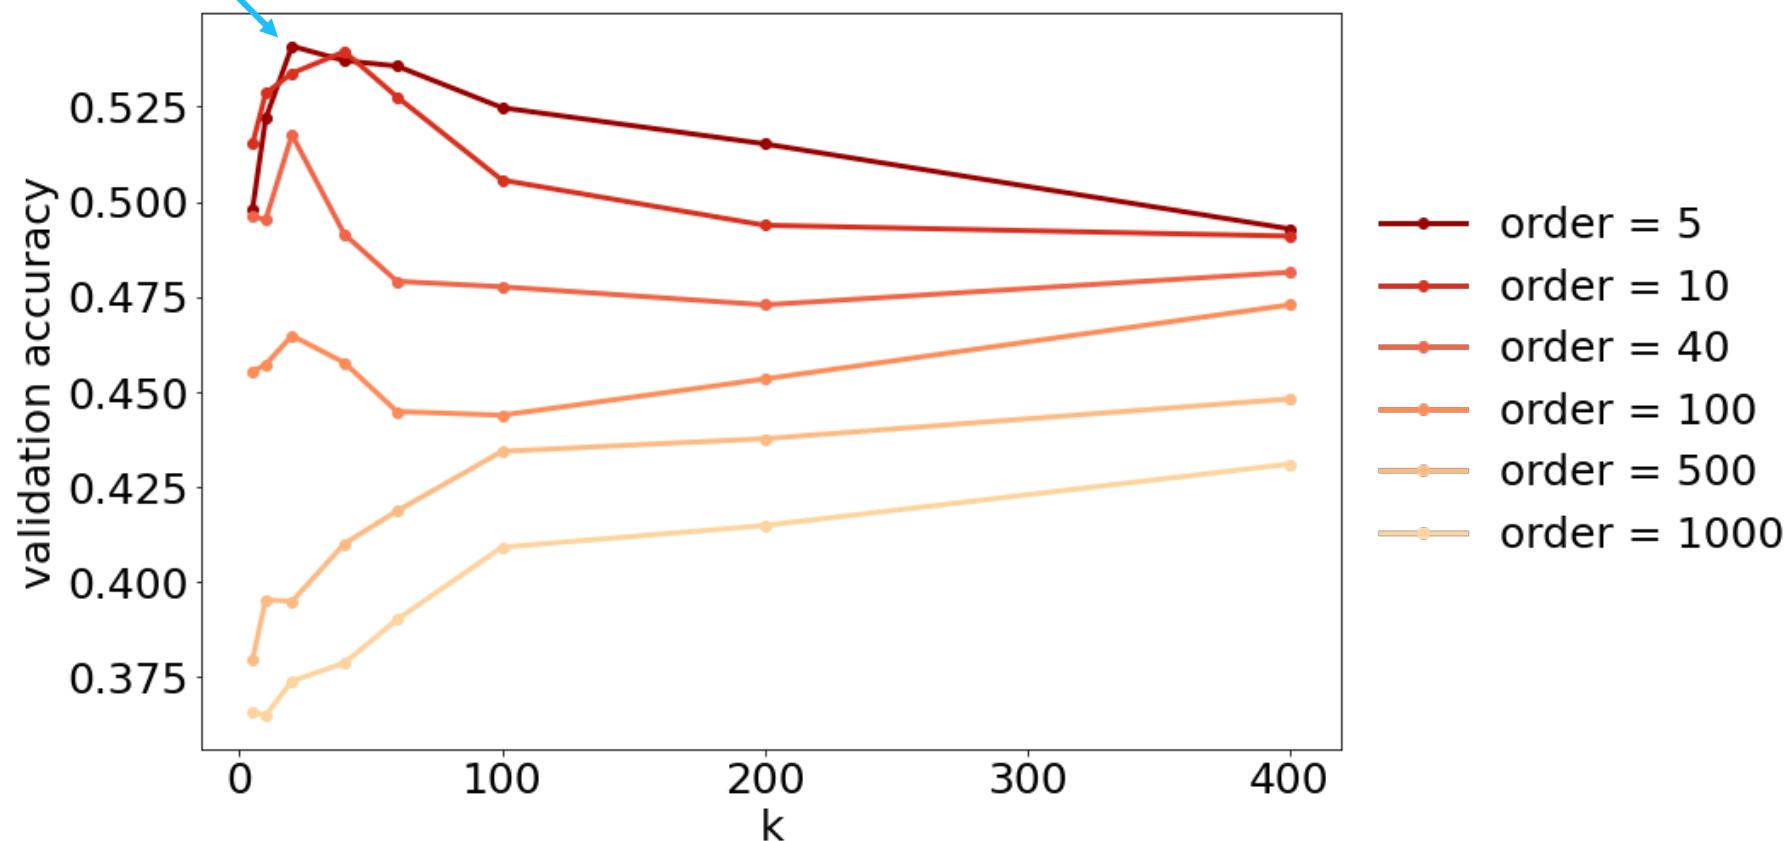
# k-NN: Hyperparameter tuning & test accuracy

Optimal parameters:  
 $r = 5, k = 20$



Testing accuracy =

51.3%



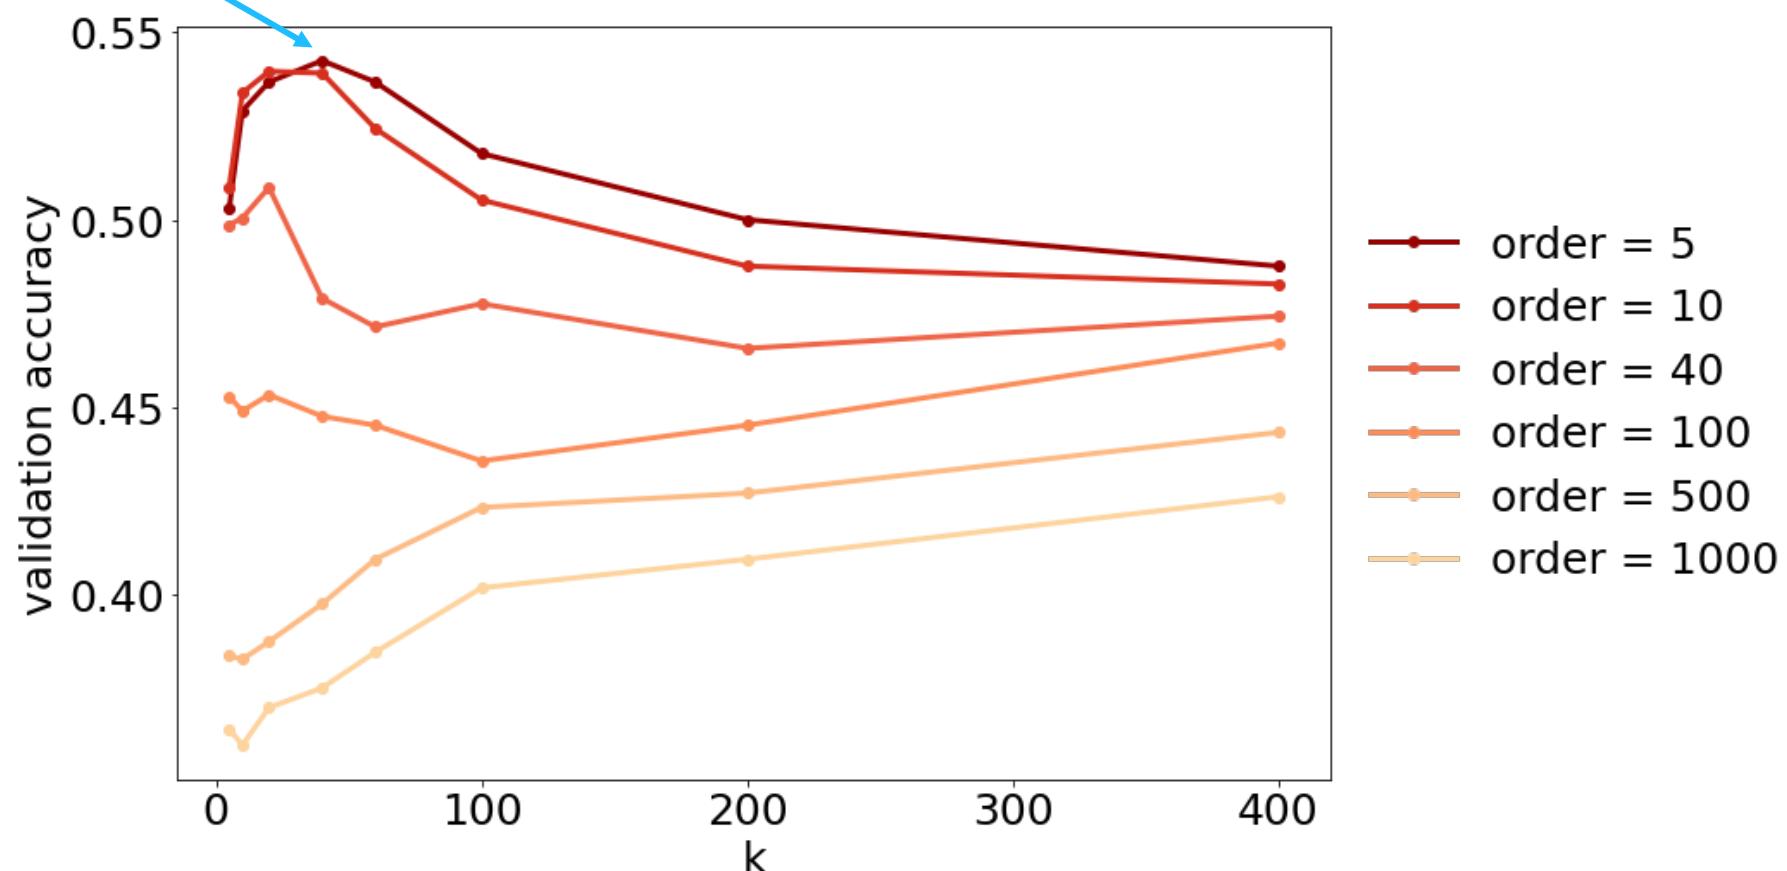
# k-NN: Add centering/scaling

Optimal parameters:  
 $r = 5, k = 40$



Testing accuracy =

53.3%



# SVM: Problem setup

**Goal:** To find an optimal separation boundary (**a hyperplane**) between two classes such that the distance between the **support vectors** of the classes (**the margin**) is maximized.

Primal form:  $\min_{\mathbf{w}, b, \xi}$

$$\|\mathbf{w}\|_2^2 + C \sum_{j=1}^n \xi_j$$

such that  $y_j \cdot (\langle \mathbf{w}, \mathbf{x}_j \rangle - b) \geq 1 - \xi_j, \quad \xi_j \geq 0$

Allows for misclassification

Weights vector normal  
to the hyperplane

## Important:

- This problem is formulated as binary classification.
- SVM is not scale invariant, therefore data needs to be standardized before training the classifier.

Rewriting the optimization problem in its **dual form** through Lagrange multipliers, allows us to use the **kernel trick** to include **nonlinearity in the hyperplane**

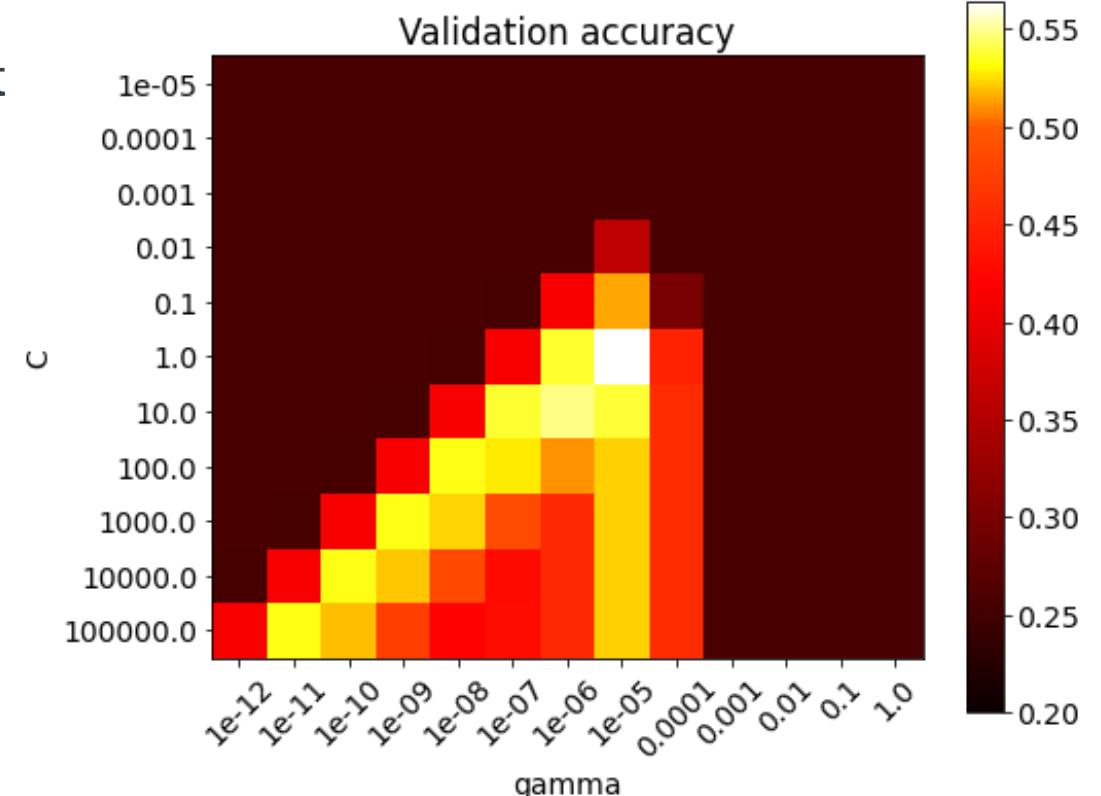
Linear:  $\langle x, x' \rangle$

<b>Evaluated kernels:</b>	Polynomial: $(\gamma \langle x, x' \rangle + r)^d$
	Sigmoid: $\tanh(\gamma \langle x, x' \rangle + r)$
	RBF: $\exp(\gamma \ x - x'\ )$

# SVM: Implementation

- Trained with enough PCA modes to capture 95% of the energy in the training dataset: **1,061 features**.
- **Data was standardized** to have zero mean and unit variance.
- Multi-class classification implemented through **one-vs-one approach**.
- **Hyperparameter tuning** strategy:
  - Logarithmic grid to assess parameters spanning multiple orders of magnitude.
  - Training and validation datasets were combined and split into validation and training portions in 5 different ways for cross validation.
  - Optimal set of parameters obtained averaging the parameters yielding the highest accuracy for each split.

**Heat map of hyperparameters logarithmic grid for RBF kernel**



# SVM: Results

---

Training with the **reduced feature space**

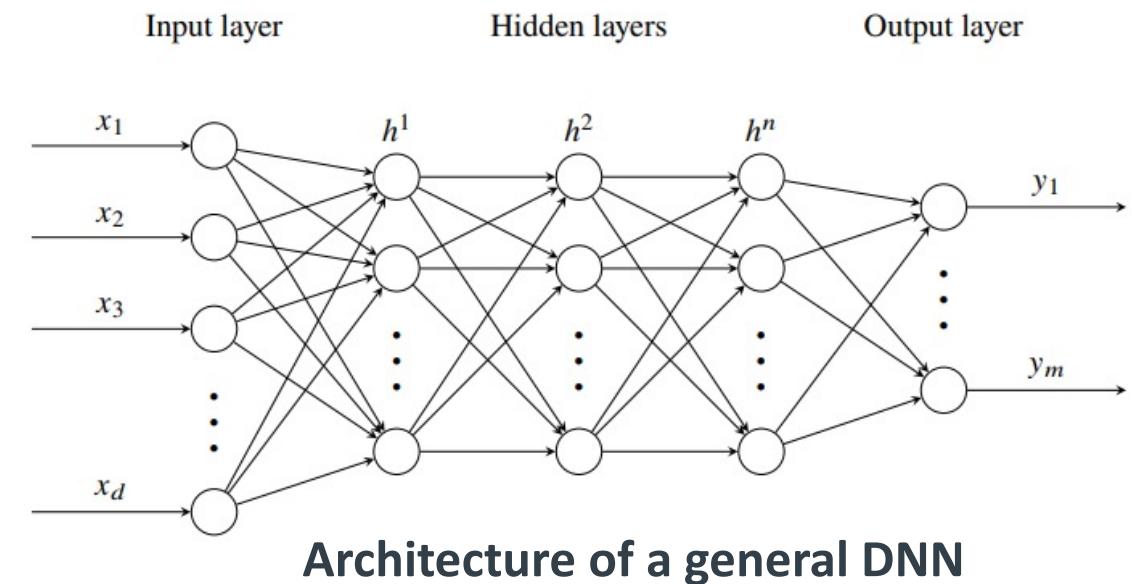
Kernel	Accuracy	
	Training	Test
Polynomial (degree 3)	98%	47%
Linear	66%	54%
Sigmoid	66%	54%
RBF	77%	56%

Training with the **raw feature space** using the same optimal hyperparameters

Kernel	Accuracy	
	Training	Test
Polynomial (degree 3)	100%	47%
Linear	68%	54%
Sigmoid	68%	54%
RBF	80%	56%

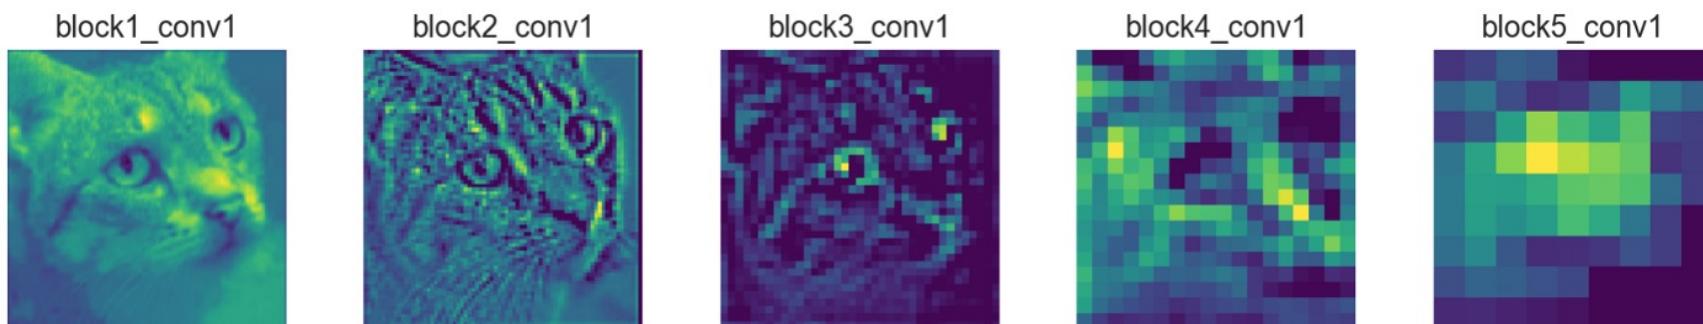
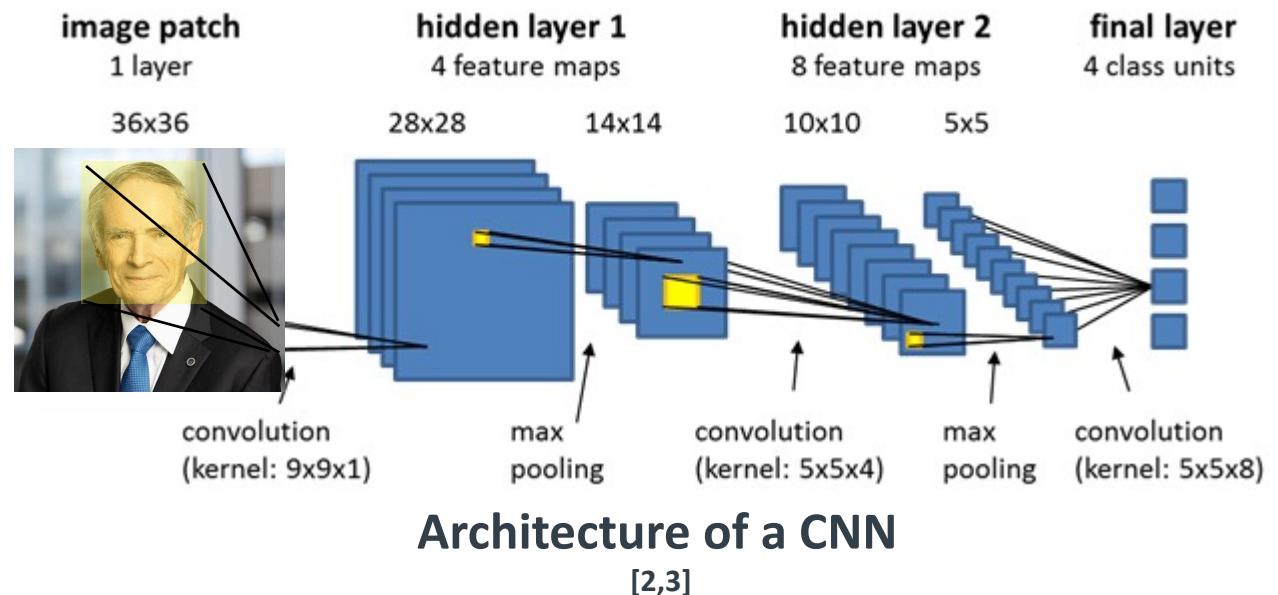
# Deep Neural Networks (DNN)

- DNNs train neurons with weights and non-linear activation functions, grouped in layers, to minimize a loss function and predict on new data
- DNNs inherently scale poorly to highly dimensional data
  - e.g., a  $256 \times 256$  pixel resolution image would require 65536 weights per neuron in dense layer
- We use a class of DNNs, referred to as convolutional neural networks (CNNs)



# Convolutional Neural Networks (CNNs)

- Convolutional layers:
  - extract features and learn them at the same time
  - use significantly less parameters than a fully connected layer
- Input: (# of images) x (image height) x (image width) x (image channels)
- Output: (# of images) x (map height) x (map width) x (map channels)

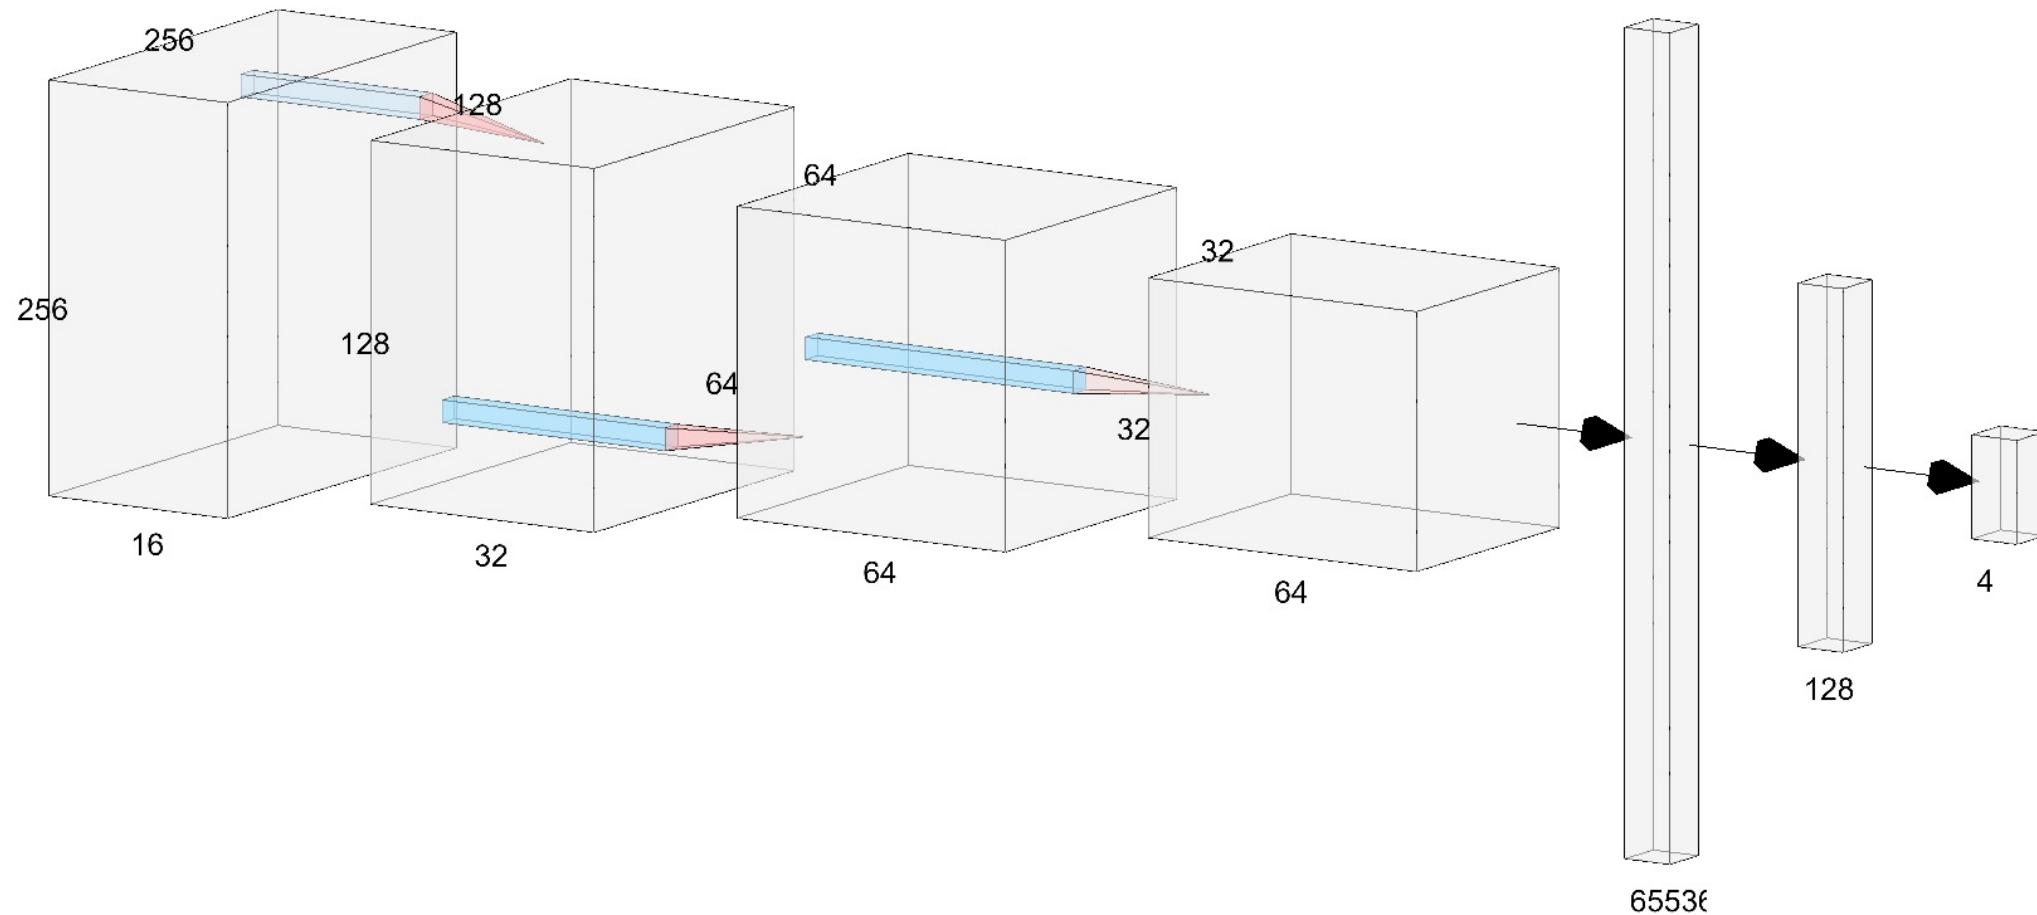


One feature map per layer in a CNN

[4]

# CNN model architecture

---



# CNN implementation

---

## Training

- Data preprocessing:
  - Each batch had rescaling, a random chance of vertical flips, rotations, shears, width shifts, height shifts, zooming, and horizontal flips
- Optimizer: Adam ( $\alpha = 1e-3$ )
- Loss: Categorical cross-entropy
- Callbacks:
  - Learning rate scheduler
  - Early stopping

## Network

- Activation Function: ReLU
- Dropout layers (at 20%)
- L2 Kernel Regularization ( $\lambda = 1e-3$ )

# CNN in the reduced space

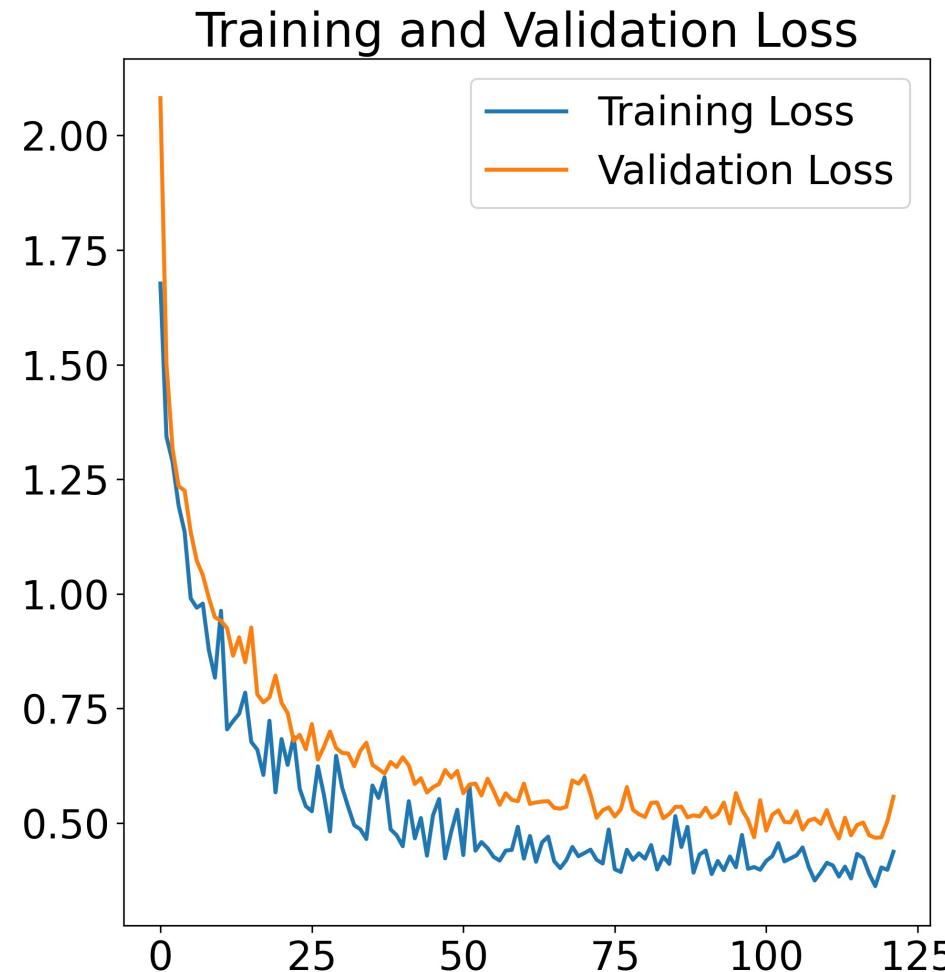
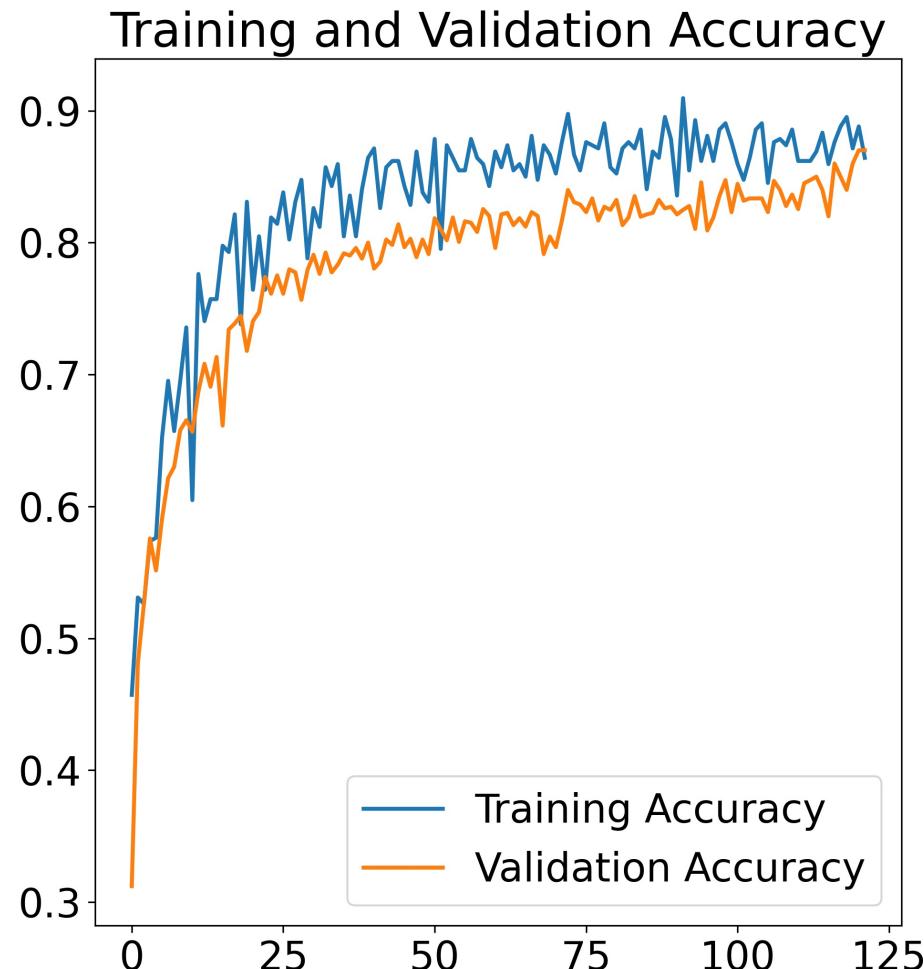
---

- Apply a CNN to projected inputs via PCA
- Attempt same network architecture, with new hyperparameters
- Experiment with variants that fit new data
- Use similar implementation tools as the CNN in full dimensional space

# CNN training history

---

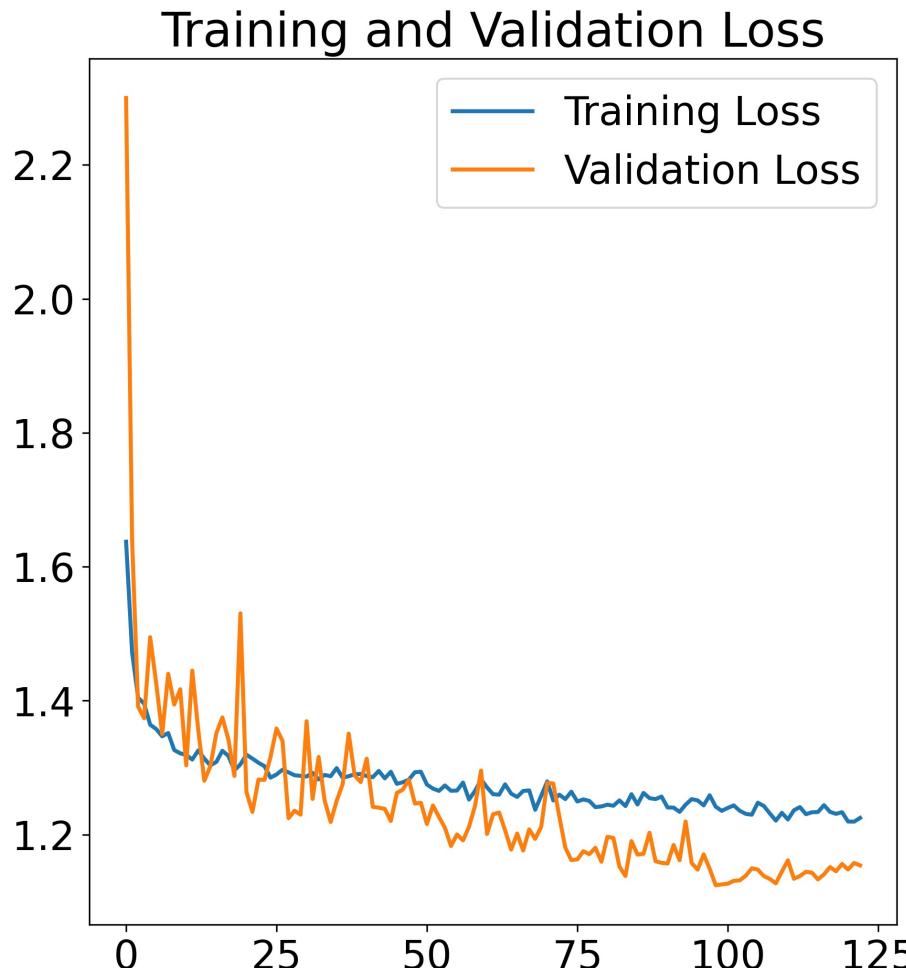
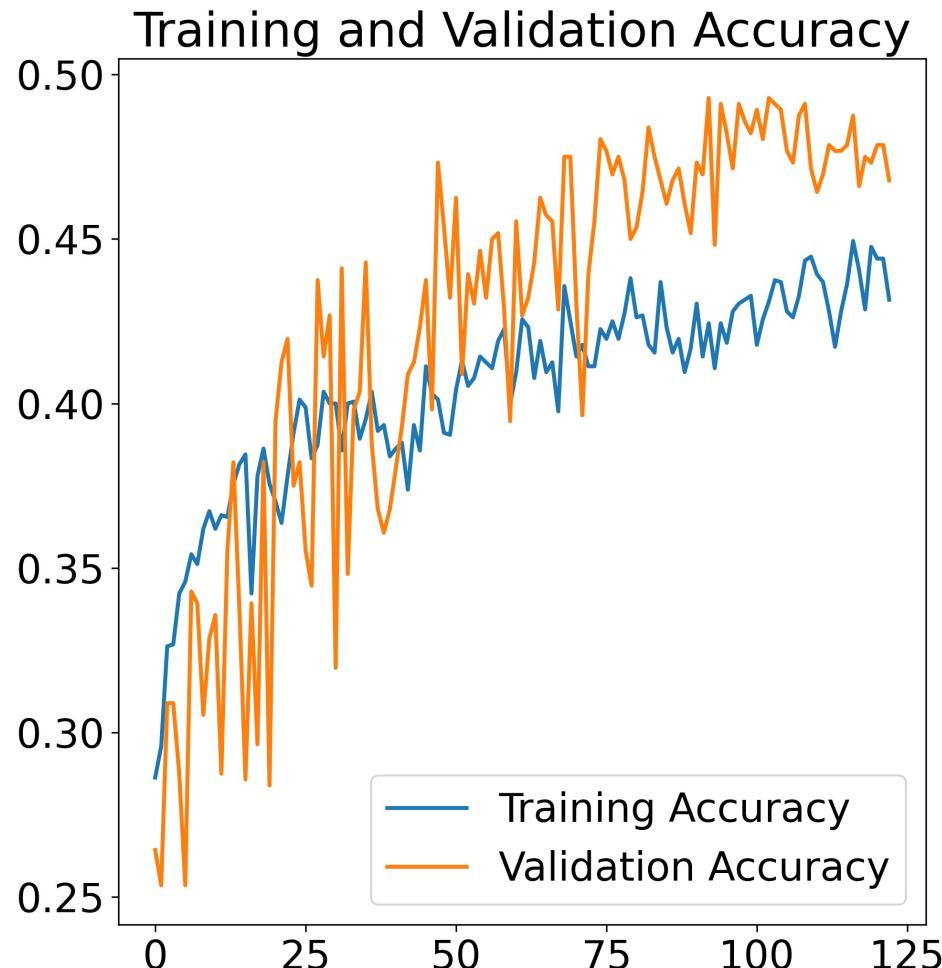
Testing accuracy = 87.0%



# CNN training history in reduced space

---

Testing accuracy = 47.3%



# Applying PCA in context

---

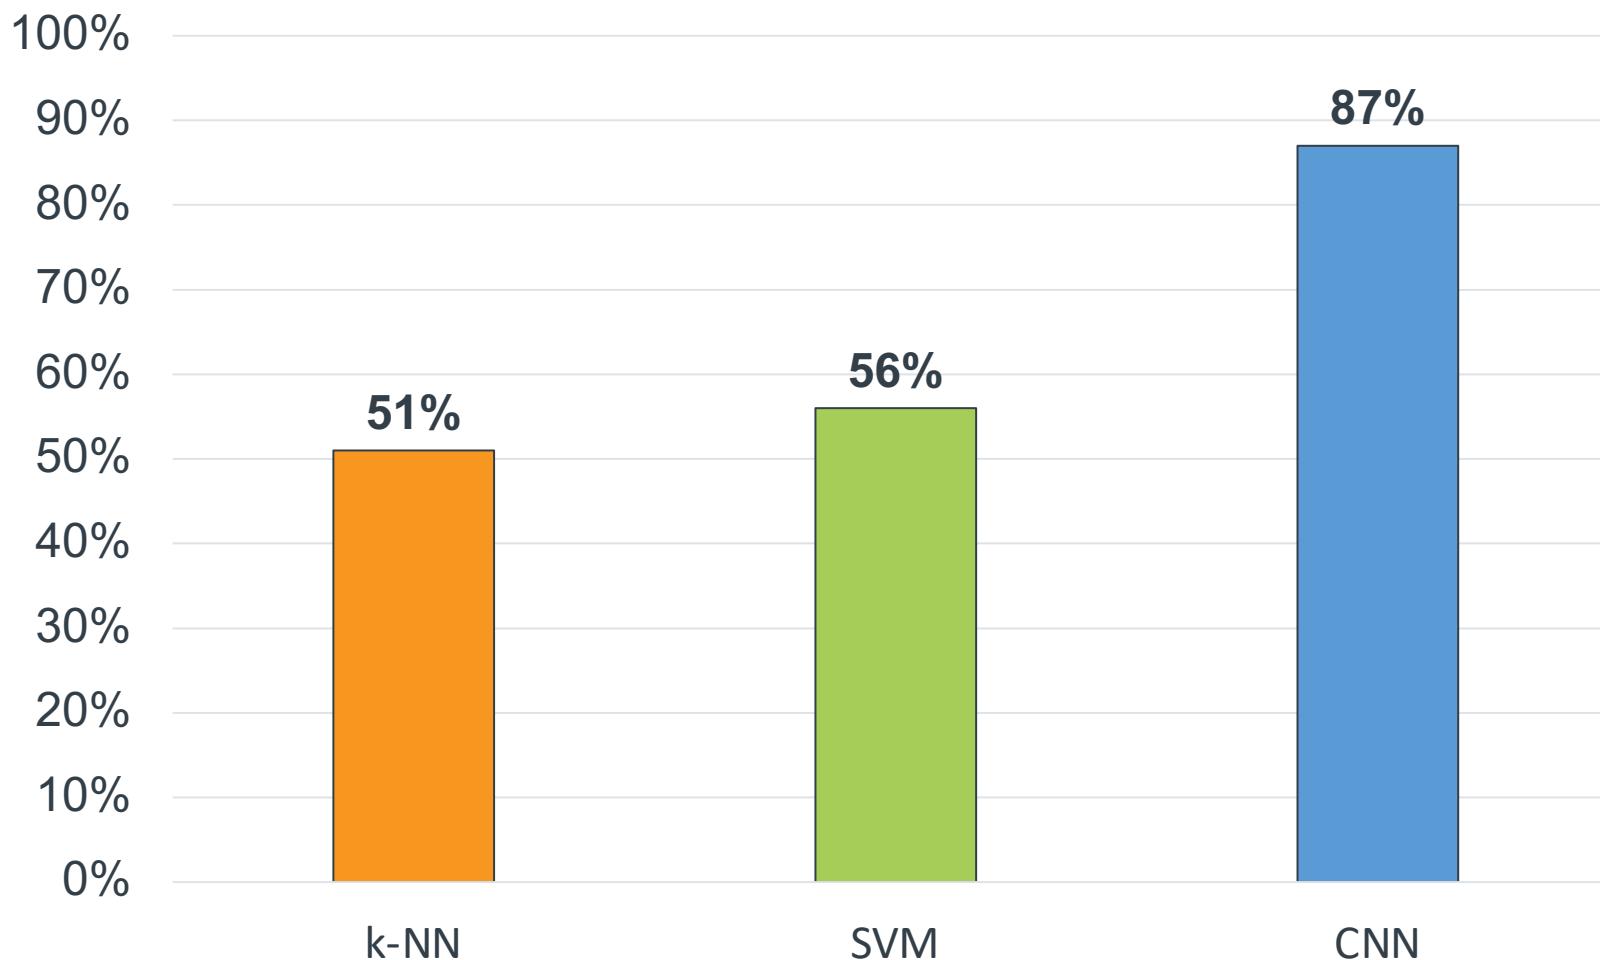
## 1. PCA ill-advised in context,

1. PCA has implications on the correlations and dimensionality of the data
2. CNNs are well equipped to leverage high dimensional data and highly correlated variables

## 2. CNN unable to generalize to test data, given training on training feature space/principal components

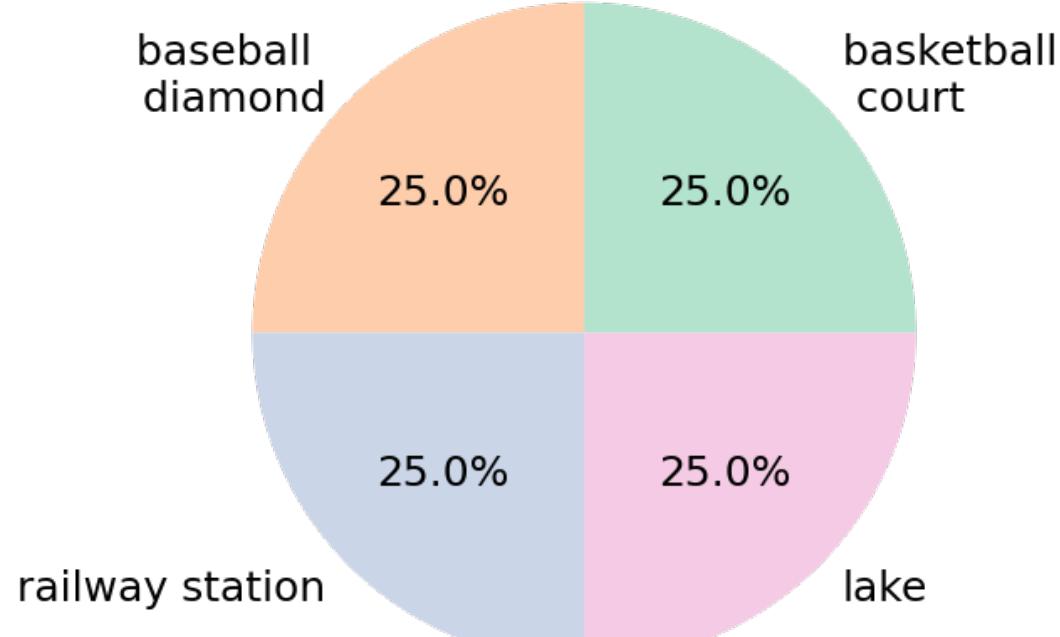
# Comparing the test accuracy

---

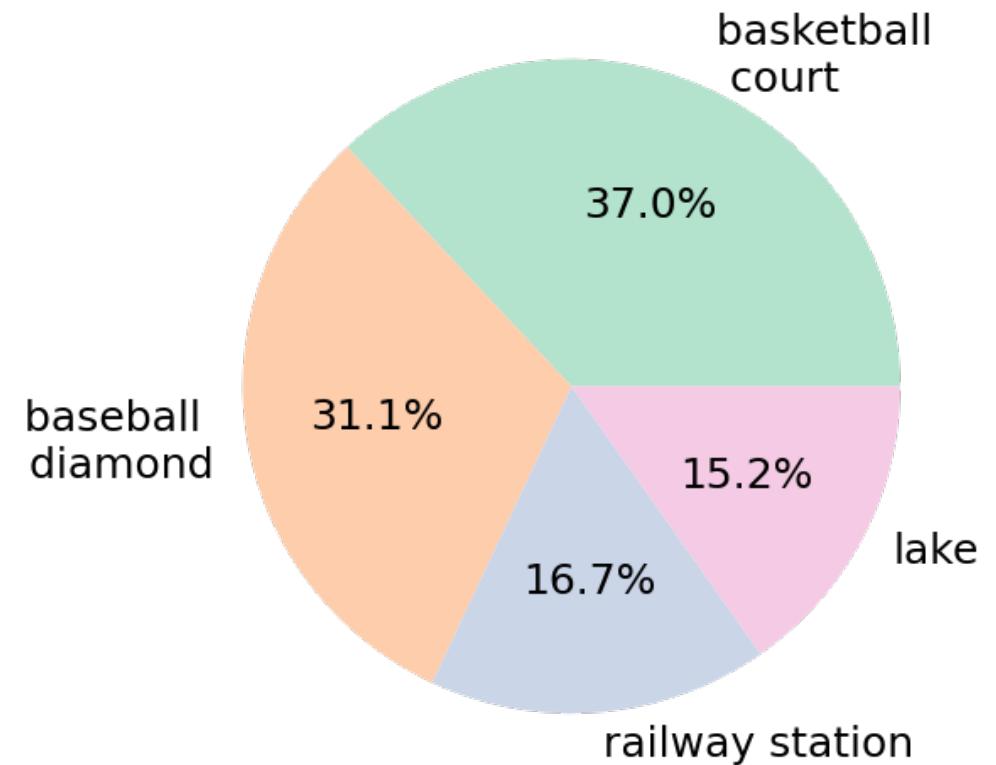


# k-NN Misclassifications

---



Training/Testing datasets Composition

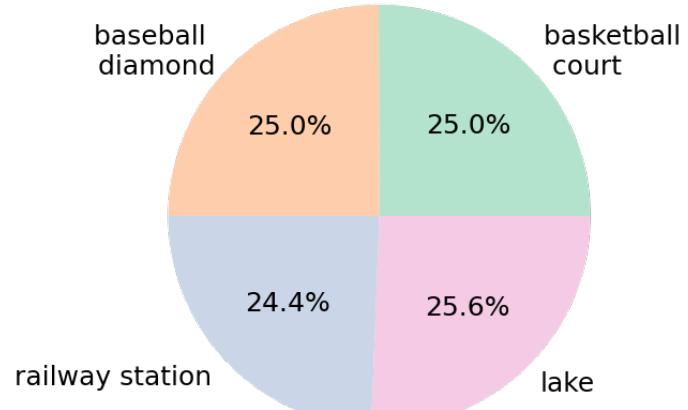


Percent of misclassifications  
corresponding to each class

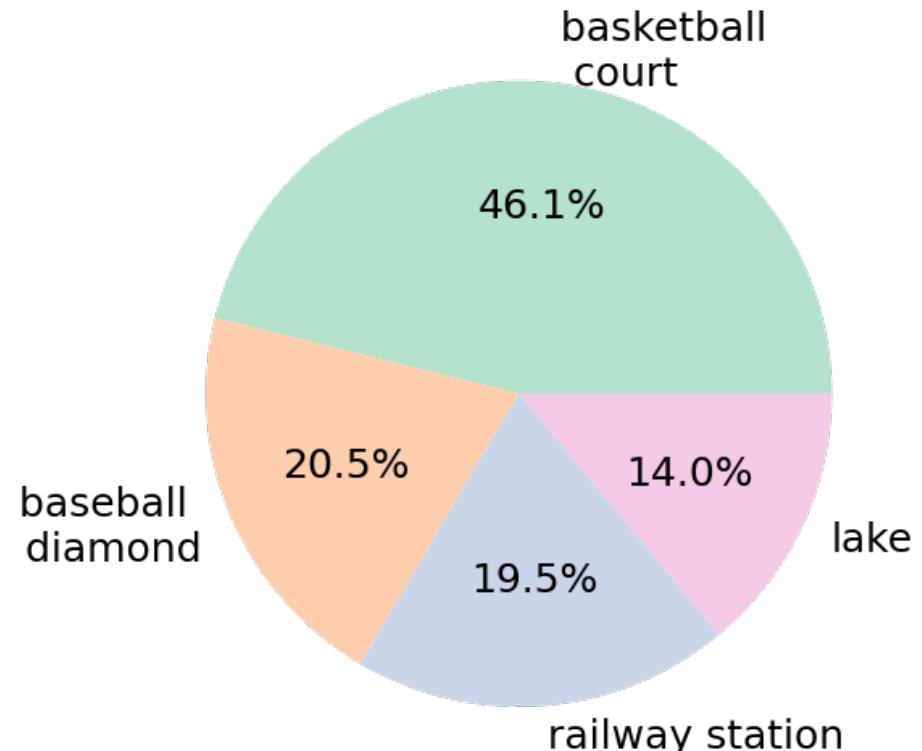
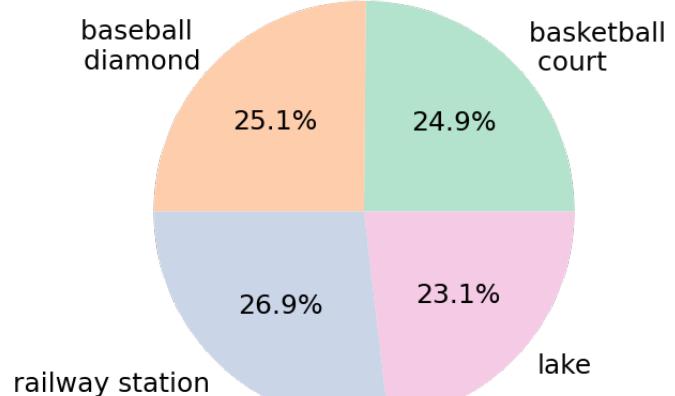
# SVM Misclassifications

## Dataset composition

Train dataset



Test dataset

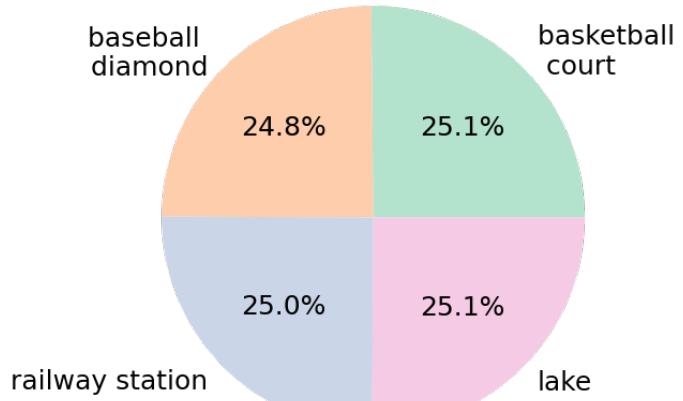


Percent of misclassifications  
corresponding to each class

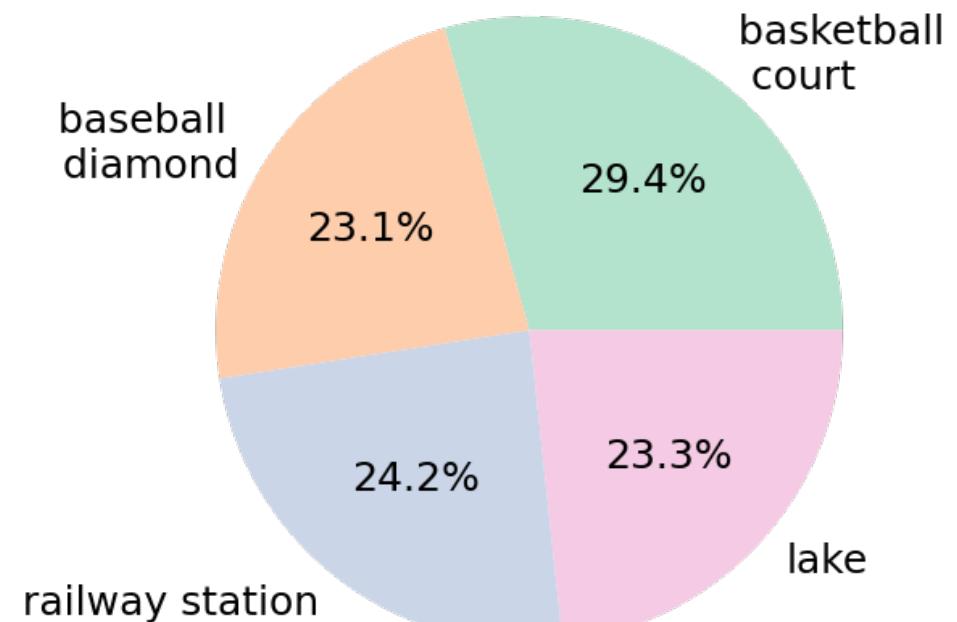
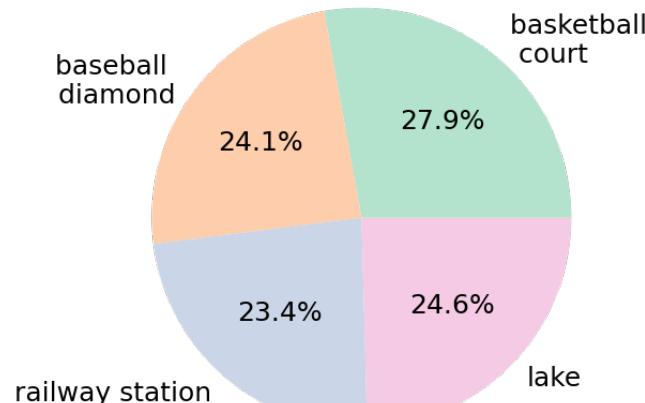
# CNN Misclassifications

## Dataset composition

Train dataset



Test dataset

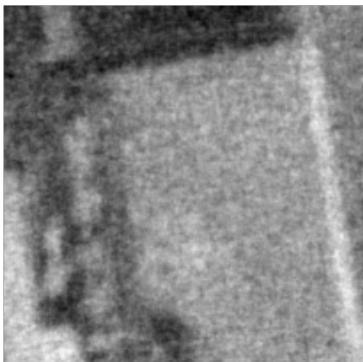


Percent of misclassifications  
corresponding to each class

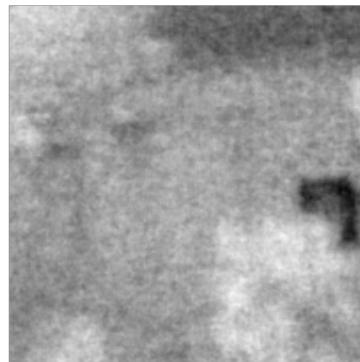
# Let's try classifying some images ourselves

---

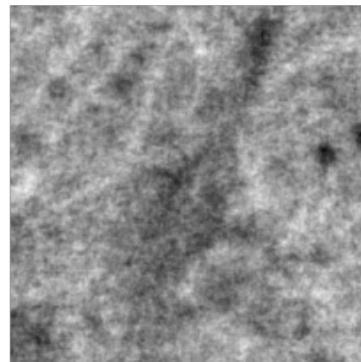
Possible classes: Basketball court, baseball diamond, railway station, lake



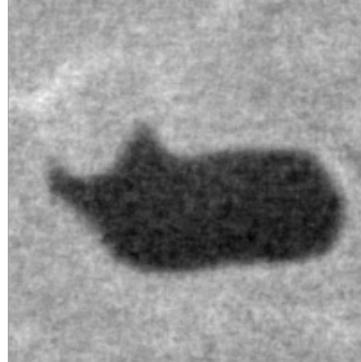
Basketball  
court



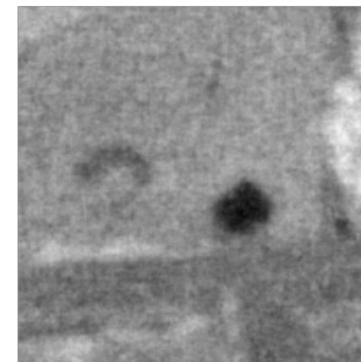
Basketball  
court



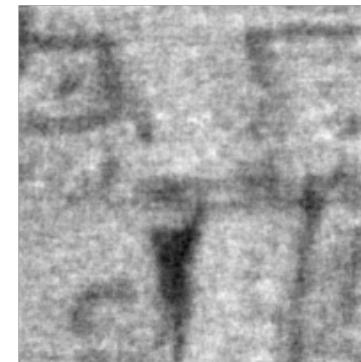
Railway  
station



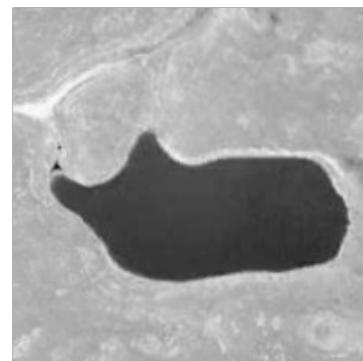
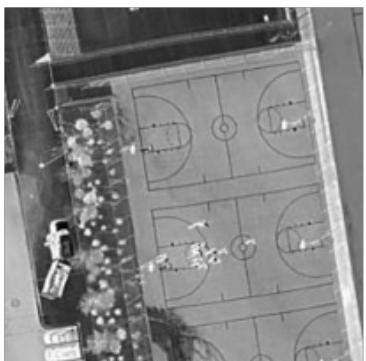
Lake?



Baseball  
diamond



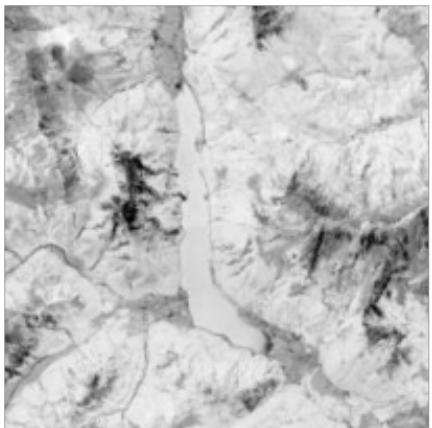
Baseball  
diamond



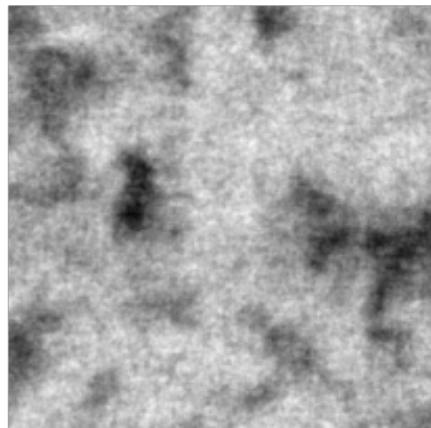
# Let's look at some lakes a little closer

---

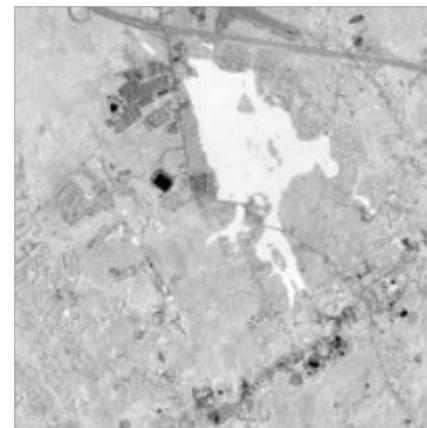
Original



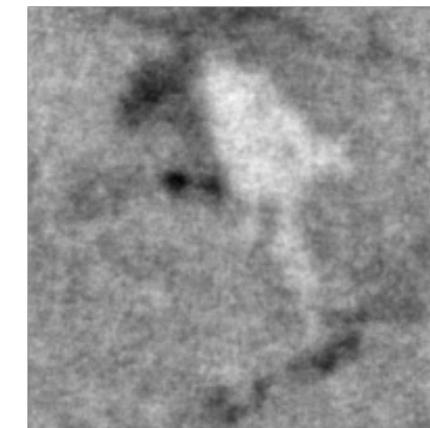
Reduced



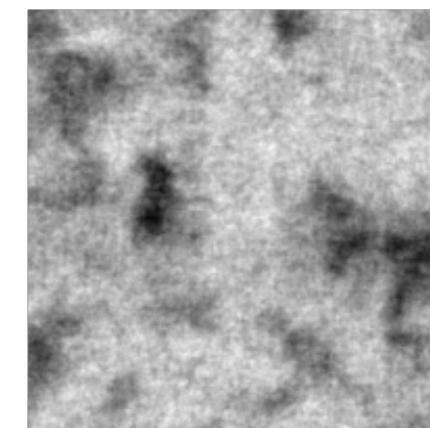
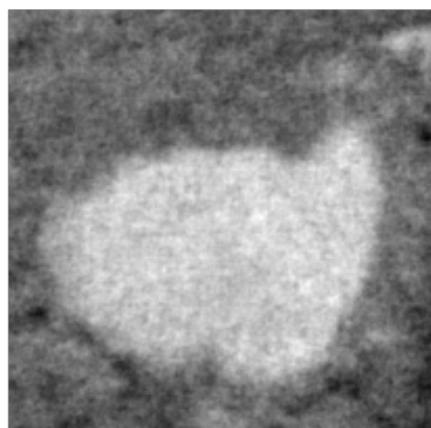
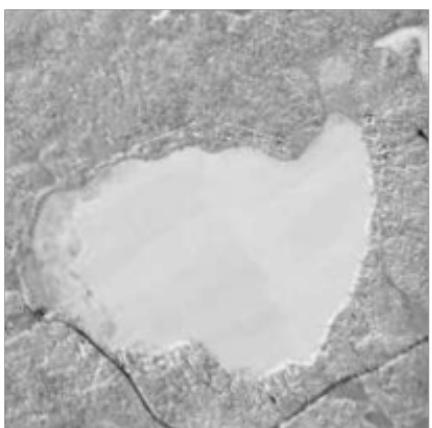
Original



Reduced



The features defining a lake are generally stronger than for the rest of classes!



# References

---

1. G. Cheng, J. Han and X. Lu, "Remote Sensing Image Scene Classification: Benchmark and State of the Art," in Proceedings of the IEEE, vol. 105, no. 10, pp. 1865-1883, Oct. 2017,  
doi: [10.1109/JPROC.2017.2675998](https://doi.org/10.1109/JPROC.2017.2675998)
2. <https://www.oden.utexas.edu/media/directory-assets/profile-imgs/oden.jpg>
3. [https://docs.ecognition.com/v9.5.0/Resources/Images/ECogUsr/UG\\_CNN\\_scheme.png](https://docs.ecognition.com/v9.5.0/Resources/Images/ECogUsr/UG_CNN_scheme.png)
4. [https://miro.medium.com/max/1400/1\\*OuxhgVj1WDDfo5UO5GIhgA@2x.png](https://miro.medium.com/max/1400/1*OuxhgVj1WDDfo5UO5GIhgA@2x.png)



The University of Texas at Austin  
**Oden Institute for Computational  
Engineering and Sciences**